

# Self-Healing Microservices with Kubernetes

Self-Adaptation in Micro-Service Architectures with Kubernetes Seminar – Summer Term 2019

Sebastian Schmidl<sup>1</sup>

**Abstract:** Abstract goes here.

**Keywords:** Self-Adaptive Systems; Self-Healing; Microservices; Cloud Computing; Kubernetes; Decentralized; Distributed; Orchestration

## 1 Introduction

1. cloud
  - cloud computing de-facto standard in industry
  - reasons: more flexibility, higher and dynamically available performance, and competitive prices [To15]
  - more hardware means more hardware can fail → plan for failure [Ne15]
  - need for resilient systems [Bo14]
  - achieve it via replication, containment, isolation, and monitoring paired with responsive actions to failures
2. microservices
  - way to allow scaling of applications combined with a way to realize containment and isolation on business boundaries
  - decompose software application into small, lightweight, autonomous services (= scaling units)
  - embrace failure: services relying on other services should deal with them failing [Ne15]
3. deployment and orchestration → Kubernetes
4. self-adaptive systems
  - self-\* properties
  - MAPE-K loop
5. self-healing

---

<sup>1</sup> Hasso Plattner Institut, University of Potsdam, Prof.-Dr.-Helmert-Str. 2-3, 14482 Potsdam, sebastian.schmidl@student.hpi.de

- currently widely-used definition for self-healing systems is from Ghosh et al. [Gh07]:

The key focus [...] is that a self-healing system should recover from the abnormal (or “unhealthy”) state and return to the normative (“healthy”) state, and function as it was prior to disruption.

- Neither fault-tolerant systems, nor survivable systems include recovery oriented functionalities that bring the system back to the healthy state, which is the key aspect of self-healing systems [Gh07].
- Combination of [PD11]
  - Fault-tolerant (handle transient failures and mask permanent ones)
  - self-stabilizing (non-fault masking; system converges to legal state in finite time and tries to remain in the same (closure))
  - survivable (maintain essential service and recover non-essential after intrusions have been dealt with)

#### 6. self-healing of microservices in cloud environments

- Psaiar; Dustdar compare self-healing in cloud environments to the techniques that achieve continuous availability of the application [PD11]:

In cloud environments, self-healing can be considered as the techniques to achieve continuous availability, which involves detecting disruptions, diagnosing failures and recovering with a sound strategy.

- In a cloud environment and a VM or container deployment, all failures are reduced to a single one: service unavailable after a while

## 2 Related Work?

- Reactive Manifesto [Bo14] asks for more resilient and responsive systems. The resilience is achieved by replication, containment, isolation, and delegation. Recovery should be handled by an external component. This could be a self healing component.
- [To15]
- [St17]
- [FN16]
- [DHT02]
- Kubernetes and alternatives

### 2.1 Self-Healing

1. sub control loop (Detect – Analyze – Recover)
2. different levels of self-healing (architecture-based, model-based, hierarchical, etc.)
3. self-healing management logic external and internal to the managed application  
**external to application**

- self-healing and management logic is run in isolation from the application code
- Examples: using services from the infrastructure provider, using third party services, or building an ad-hoc solution (e. g. using Kubernetes) [To15]
- current state of the art for monitoring, health management, and scaling logic
- could lead to vendor lock-in
- external management logic has to be themselves resilient, fault-tolerant, and scalable

#### **within application**

- approach by Toffetti et al. for microservices; leverages standard methods from distributed systems (such as consensus algorithms) to assign self-management functionality to nodes of the application; hierarchical approach [To15]

## **2.2 Kubernetes**

1. what is Kubernetes?
2. architecture and how it works

## **3 Using Kubernetes to implement a self-healing application**

1. How would a setup of a self-healing microservice architecture look like?
2. self-healing properties available in Kubernetes
 

A Controller can create and manage multiple Pods for you, handling replication and rollout, and providing self-healing capabilities at cluster scope.

  - recovery of stateful applications:
    - Deployment definition via StatefulSet: <https://kubernetes.io/docs/tutorials/stateful-application/basic-stateful-set/>
    - Uses PersistentVolumes (provided by the underlying cloud platform, e.g AWS, GCP, OpenStack) for storage
    - Pods have a unique identity (name, network id, K8s configuration)
    - Failed pods will be rescheduled on other nodes with their identity (re-using the assigned persistent volume and network id)
    - A headless Service takes care of service discovery using SRV records and DNS (re-routing traffic to rescheduled pods on different nodes)
    - therefore, relies on the availability and fault-tolerance of the used persistent volumes
  - recovery of stateless applications:
    - Deployment definition via Deployment and the specification of replicas > 1 or with ReplicaSet

- Failing pods will be recreated to match the desired number of replicas (node placement is transparent)
- daemons: applications per node
  - Defined via DaemonSets: <https://kubernetes.io/docs/concepts/workloads/controllers/daemonset>
  - Ensures (monitors, restarts) that a copy of an application is run on each node (also on adding or removing nodes)
- 3. additional configuration and tools needed

## 4 Discussion

1. limitations
  - **external management logic has to be themselves resilient, fault-tolerant, and scalable**
  - Kubernetes default only one master → HA setup across availability zones
  - quite a lot of configuration work, not automation yet (WIP)
  - only one master will be active (the other two will be passive), full state replication via etcd
  - fail-over will be handled by load balancer component
  - **only external view on the system**
2. benefits
3. interesting facts and insights

## 5 Conclusion

### References

- [Bo14] Bonér, J.; Farley, D.; Kuhn, R.; Thompson, M.: The Reactive Manifesto, 2014, URL: <https://www.reactivemaneifesto.org>, visited on: 06/04/2019.
- [DHT02] Dashofy, E. M.; van der Hoek, A.; Taylor, R. N.: Towards Architecture-based Self-healing Systems. In: Proceedings of the First Workshop on Self-healing Systems (WOSS). Pp. 21–26, 2002.
- [FN16] Florio, L.; Nitto, E. D.: Gru: An Approach to Introduce Decentralized Autonomic Behavior in Microservices Architectures. In: IEEE International Conference on Autonomic Computing (ICAC). Pp. 357–362, 2016.
- [Gh07] Ghosh, D.; Sharman, R.; Raghav Rao, H.; Upadhyaya, S.: Self-healing Systems - Survey and Synthesis. Decision Support Systems 42/4, pp. 2164–2185, 2007.
- [Ne15] Newman, S.: Building Microservices: Designing Fine-Grained Systems. O'Reilly Media, 2015.

- [PD11] Psailer, H.; Dustdar, S.: A survey on self-healing systems: approaches and systems. *Computing* 91/1, pp. 43–73, 2011.
- [St17] Stack, P.; Xiong, H.; Mersel, D.; Makhloufi, M.; Terpend, G.; Dong, D.: Self-Healing in a Decentralised Cloud Management System. In: *Proceedings of the 1st International Workshop on Next Generation of Cloud Architectures*. 3:1–3:6, 2017.
- [To15] Toffetti, G.; Brunner, S.; Blöchliger, M.; Dudouet, F.; Edmonds, A.: An architecture for self-managing microservices. In: *Proceedings of the 1st International Workshop on Automated Incident Management in Cloud*. Pp. 19–24, 2015.